



Mark Lemmert &lt;mark.lemmert@gmail.com&gt;

---

## driver file read question

---

Peter Ferrie &lt;peter.ferrie@gmail.com&gt;

Mon, Aug 8, 2016 at 5:37 PM

To: Mark Lemmert &lt;mark.lemmert@gmail.com&gt;

Hi Mark,

>> It required a fairly substantial change internally, owing to the way  
>> that the code grew over time. Initially, it was based on a  
>> sector-oriented DOS 3.3 driver. Then it became a sector-oriented  
>> ProDOS driver, and has finally become a block-oriented ProDOS driver.  
>  
> Ah, ok. Sounds like efficiency probably increased with each step as the  
> design became closer to the format of the raw data?

It's become simpler to read in blocks, but more complicated to allow  
reads to resume, because of the need to keep a pointer into a cache of  
the last read block.  
Still, it will be faster than ProDOS.

>> Since I'm making significant changes right now, is there any chance  
>> that you'd want a seek routine? Or will you always read initially  
>> from the start of the file and discard what isn't interesting?  
>  
> I have some ideas on how that might be useful. Thanks! A few questions.  
>  
> With a seek routine would that mean I could specify a range of bytes in the  
> file I want to read in? Like read byte \$100-\$180, skipping the first page  
> of data? Or would it be more like DOS RWTS where I'd specify the block  
> range to read and the driver returns whatever is in those blocks?

The seek would be a separate command and a separate call from the  
read. You specify the number of bytes to skip in one call, and the  
number of bytes to read from the new position using another call.

> Is the advantage of seek routine speed and/or simplicity in the higher  
> level code? For example, if the driver does a read/seek to return bytes  
> \$5000-\$10000 do you think the driver run time would be significantly faster  
> than if the driver just read bytes \$00 - \$10000 and let the higher level  
> code discard bytes \$00-\$4FFF?

Yes, the seek reads only one block (the final one where the file  
pointer ends up, so that the next read comes from the cache  
initially), instead of reading and decoding everything in between.  
That's a huge saving if you're seeking a long way.

> Would an upgrade the driver be available first with just the read length  
> added, so I can start working with it, or does the seek need to be done at  
> the same time?

I've added the whole thing. It's ready to use right now. You have to  
specify the size value for all of the commands now, not just the  
write.

Some notes about it:

- ;- writes can be performed beginning at any block in a file, by seeking to the block first
- ;- file size can be retrieved by seeking to position 0 within a file and then reading bleft
- ;- to read a complete file, set the size to something large, eg \$ffff

;constants

cmdseek = 0

cmdread = 1

cmdwrite = 2

;zpage usage

tmpsec = \$3c

reqsec = \$3d

A1L = \$3c ;only during init

A1H = \$3d ;only during init

A2L = \$3e ;only during init

A2H = \$3f ;only during init

A3L = \$40 ;only during init

A3H = \$41 ;only during init

curtrk = \$40

command = \$42 ;ProDOS constant

unit = \$43 ;ProDOS constant

adrlo = \$44 ;ProDOS constant

adrhi = \$45 ;ProDOS constant

bloklo = \$46 ;ProDOS constant

blokhi = \$47 ;ProDOS constant

bleftlo = \$ef ;(internal) bytes left in file

blefthi = \$f0 ;(internal) bytes left in file

blkofflo = \$f1 ;(internal) offset within cache block

blkoffhi = \$f2 ;(internal) offset within cache block

status = \$f3 ;returns non-zero on error

auxreq = \$f4 ;set to 1 to read/write aux memory

sizelo = \$f5 ;size of request

sizehi = \$f6 ;size of request

entries = \$f7 ;(internal) total number of entries

reqcmd = \$f8 ;0=seek, 1=read, 2=write

ldrlo = \$f9 ;set to load address

ldrhi = \$fa ;set to load address

namlo = \$fb ;set to name of file to open

namhi = \$fc ;set to name of file to open

step = \$fd ;(internal) state for stepper motor

tmptrk = \$fe ;(internal) temporary copy of

current track

phase = \$ff ;(internal) current phase for seek

reloc = \$d000

dirbuf = reloc+\$500 ;\$200 bytes

encbuf = dirbuf+\$200 ;\$200 bytes

Memory usage is now \$D000-D8FF.

The openfile entrypoint is \$D003. It supports seek/read/write as well.

The general seek/read/write entrypoint is \$D000 (open the file first).

There is no "close" operation. You can have only one file open at a time, so opening another file will close the current one.

;open file and read 3 bytes to \$1234

```
ldx #cmdread
stx reqcmd
dex
stx auxreq
lda #3
sta sizelo
lda #0
sta sizehi
lda #$34
sta ldrlo
lda #$12
sta ldrhi
lda #<noxmain
sta namlo
lda #>noxmain
sta namhi
jsr $d003
```

;seek to position \$192 in opened file

```
lda #cmdseek
sta reqcmd
lda #$92
sta sizelo
lda #1
sta sizehi
jsr $d000
```

;read \$500 bytes to \$6000 from opened file

```
ldx #cmdread
stx reqcmd
lda #0
sta sizelo
lda #5
sta sizehi
lda #0
sta ldrlo
lda #$60
sta ldrhi
jsr $d000
```

Entries no longer returns the number of pages read, but you can track it yourself since you know exactly how many bytes you are reading, and you have access to the file size.

Another thing - file entries in the file system aren't "free". There are 13 files per block. The 14th file will require one more block on the disk. For each file that is larger than a block, there will be a block that holds the block list for that file. If you have, say, three files that are 513 bytes long, that's six blocks on the disk (three block lists, three data blocks).

With the new driver, you can pack many files into a single one if you want to. It could save a lot of space. We can talk more about that if it's interesting to you.

---



**NOXARCH.zip**

2K